

# S-CSIS311\_EA3\_Regression\_Imperial

September 11, 2024

*De La Salle University – Dasmariñas*  
*College of Information and Computer Studies*

*S-CSIS311 / S-CSIS311LA*  
*Introduction to Machine Learning*

*Enabling Assessment: Regression*  
*Wednesday, September 11, 2024*

**Luis Anton P. Imperial**  
**BCS32**

## 1 Enabling Assessment – Regression

### 1.1 Learning Outcome:

1. Students will be able to explain the basic concept of linear regression.
2. Students will be able to build a simple regression model in Google Colab.
3. Students will be able to use Google Colab to visualize key metrics.

### 1.2 Direction:

Using Anaconda or Google Colab, solve the machine problem. Evaluate, analyze, and explain the steps using your chosen tool. Use the dataset housing. After completing the solution, create two copies: one in PDF format and one in Python (.py) format. Submit both files to the folder provided in MS Teams.

### 1.3 Steps:

1. Import the libraries
2. Load the data
3. Setup the data
4. Prepare the data processing
5. Train the model
6. Test the model
7. Print the predictions
8. Visualize the model

### 1.4 Objective:

To predict the Median House Value based on Median Income

```
[1]: # Importing libraries required
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import datasets, linear_model
from sklearn.metrics import mean_squared_error, r2_score

SEED = 42
np.random.seed(SEED)
```

## 1.5 Questions:

1. Load the dataset
2. Load the data frame.head()
3. Load the df = df[['median\_income', 'median\_house\_value']]
4. Display dataframe information using df.info() and display df.head(10)
5. What values do you see?
6. What distributions do you see?
7. What relationships do you see?
8. What relationships do you think might benefit the prediction problem?
9. What ideas about the domain does the data spark?
10. Display all columns, head, tail and info using df.columns, df.head(), df.tail() and df.info()
11. Get an overall sense of the data shape using df.describe()
12. Group the data and sort in ascending order
13. Plot the data using scatterplot
14. Removing outliers
15. Visualize the trained model

```
[13]: ##### 1. Load the dataset
from google.colab import drive
drive.mount('/content/drive')
url = '/content/drive/MyDrive/Documents/Reference Documents/BCS3 CSIS311_
↳Reference Documents/Supplementary BCS3 CSIS311 Reference Documents/housing.
↳csv'
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True).

```
[14]: ##### 2. Load the data frame.head()
df = pd.read_csv(url)
df.head()
```

```
[14]:      longitude  latitude  housing_median_age  total_rooms  total_bedrooms  \
0      -122.23    37.88           41.0           880.0           129.0
1      -122.22    37.86           21.0          7099.0          1106.0
2      -122.24    37.85           52.0          1467.0           190.0
3      -122.25    37.85           52.0          1274.0           235.0
```

4	-122.25	37.85	52.0	1627.0	280.0
---	---------	-------	------	--------	-------

	population	households	median_income	median_house_value	ocean_proximity
0	322.0	126.0	8.3252	452600.0	NEAR BAY
1	2401.0	1138.0	8.3014	358500.0	NEAR BAY
2	496.0	177.0	7.2574	352100.0	NEAR BAY
3	558.0	219.0	5.6431	341300.0	NEAR BAY
4	565.0	259.0	3.8462	342200.0	NEAR BAY

### 1.5.1 (2/15) Loading the dataset

As we can see from the previous code block, we have properly loaded the Comma-Separated Values (CSV) file we will rely on as our dataset.

The `head()` function of the Pandas package displays the “head”, meaning the first few rows, of a given dataset. It is often used to test for the program’s access to the data.

```
[5]: ##### 3. Load the df = df[['median_income', 'median_house_value']]
df = df[['median_income', 'median_house_value']]
```

```
[6]: ##### 4. Display dataframe information using df.info() and display df.head(10)
df.info()
df.head(10)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 2 columns):
#   Column                Non-Null Count  Dtype
---  -
0   median_income          20640 non-null  float64
1   median_house_value     20640 non-null  float64
dtypes: float64(2)
memory usage: 322.6 KB
```

```
[6]:    median_income  median_house_value
0         8.3252         452600.0
1         8.3014         358500.0
2         7.2574         352100.0
3         5.6431         341300.0
4         3.8462         342200.0
5         4.0368         269700.0
6         3.6591         299200.0
7         3.1200         241400.0
8         2.0804         226700.0
9         3.6912         261100.0
```

### 1.5.2 (5/15) Peeking through people's houses

5. What values do you see?
6. What distributions do you see?
7. What relationships do you see?
8. What relationships do you think might benefit the prediction problem?
9. What ideas about the domain does the data spark?

The values provided are the median income and the median house value in each of the first ten (out of 20,640) households in the dataset.

The distribution would be somewhat of a straight line going from bottom-left to top-right, considering that a glance of the dataset's head gives the impression of a positive correlation between a neighborhood's median income and the home's median value.

Perhaps the relationships that can be formed are one-to-many and one-to-one. The median housing value can rise not just because it is located in a neighborhood with higher-income residents; it can also rise depending on other statistics such as its proximity to a body of water.

I believe one-to-many relationships that also include other factors such as the median age of the houses in each area, as well as the total number of rooms each house can provide, can help predict the value of newly built and planned houses.

In simpler terms, the data suggests that the richer the family, the more expensive their house can get. *Sounds like a fairly understood overview, isn't it?*

10. Display all columns, head, tail and info using `df.columns`, `df.head()`, `df.tail()` and `df.info()`
11. Get an overall sense of the data shape using `df.describe()`

```
[17]: ##### 10. Display all columns, head, tail and info using df.columns, df.head(),  
        ↪ df.tail() and df.info()  
df = pd.read_csv(url)  
  
df.columns  
df.head()  
df.tail()  
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 20640 entries, 0 to 20639  
Data columns (total 10 columns):  
#   Column                Non-Null Count  Dtype  
---  -  
0   longitude              20640 non-null  float64  
1   latitude               20640 non-null  float64  
2   housing_median_age     20640 non-null  float64  
3   total_rooms            20640 non-null  float64  
4   total_bedrooms         20433 non-null  float64  
5   population              20640 non-null  float64  
6   households              20640 non-null  float64  
7   median_income          20640 non-null  float64
```

```

8    median_house_value    20640 non-null    float64
9    ocean_proximity       20640 non-null    object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB

```

```
[18]: ##### 11. Get an overall sense of the data shape using df.describe()
df.describe()
```

```
[18]:
```

	longitude	latitude	housing_median_age	total_rooms	\
count	20640.000000	20640.000000	20640.000000	20640.000000	
mean	-119.569704	35.631861	28.639486	2635.763081	
std	2.003532	2.135952	12.585558	2181.615252	
min	-124.350000	32.540000	1.000000	2.000000	
25%	-121.800000	33.930000	18.000000	1447.750000	
50%	-118.490000	34.260000	29.000000	2127.000000	
75%	-118.010000	37.710000	37.000000	3148.000000	
max	-114.310000	41.950000	52.000000	39320.000000	

	total_bedrooms	population	households	median_income	\
count	20433.000000	20640.000000	20640.000000	20640.000000	
mean	537.870553	1425.476744	499.539680	3.870671	
std	421.385070	1132.462122	382.329753	1.899822	
min	1.000000	3.000000	1.000000	0.499900	
25%	296.000000	787.000000	280.000000	2.563400	
50%	435.000000	1166.000000	409.000000	3.534800	
75%	647.000000	1725.000000	605.000000	4.743250	
max	6445.000000	35682.000000	6082.000000	15.000100	

	median_house_value
count	20640.000000
mean	206855.816909
std	115395.615874
min	14999.000000
25%	119600.000000
50%	179700.000000
75%	264725.000000
max	500001.000000

```
[19]: ##### This will show only features that have nonzero missing values.
df_na = df.isna().sum()
df_na
```

```
[19]: longitude          0
latitude              0
housing_median_age    0
total_rooms           0
total_bedrooms        207
```

```

population          0
households           0
median_income        0
median_house_value   0
ocean_proximity      0
dtype: int64

```

```

[20]: ##### Limit to categorical data using df.select_dtypes()
df_cat = df.select_dtypes(include=['object'])
df_cat.nunique()

```

```

[20]: ocean_proximity      5
dtype: int64

```

```

[21]: ##### Limit to numerical data df.select_dtypes()
df_num = df.select_dtypes(include=['number'])
df_num.nunique()

```

```

[21]: longitude          844
latitude                862
housing_median_age      52
total_rooms             5926
total_bedrooms          1923
population              3888
households              1815
median_income           12928
median_house_value      3842
dtype: int64

```

```

[22]: ##### Look at correlations in the numerical independent variables as well as
↳ the dependent variables by executing df_num.corr()
df_num.corr()

```

```

[22]:
longitude  latitude  housing_median_age  total_rooms  \
longitude    1.000000 -0.924664         -0.108197    0.044568
latitude   -0.924664  1.000000          0.011173   -0.036100
housing_median_age -0.108197  0.011173         1.000000   -0.361262
total_rooms    0.044568 -0.036100        -0.361262    1.000000
total_bedrooms  0.069608 -0.066983        -0.320451    0.930380
population     0.099773 -0.108785        -0.296244    0.857126
households     0.055310 -0.071035        -0.302916    0.918484
median_income  -0.015176 -0.079809        -0.119034    0.198050
median_house_value -0.045967 -0.144160         0.105623    0.134153

total_bedrooms  population  households  median_income  \
longitude        0.069608    0.099773    0.055310   -0.015176
latitude       -0.066983   -0.108785   -0.071035   -0.079809

```

housing_median_age	-0.320451	-0.296244	-0.302916	-0.119034
total_rooms	0.930380	0.857126	0.918484	0.198050
total_bedrooms	1.000000	0.877747	0.979728	-0.007723
population	0.877747	1.000000	0.907222	0.004834
households	0.979728	0.907222	1.000000	0.013033
median_income	-0.007723	0.004834	0.013033	1.000000
median_house_value	0.049686	-0.024650	0.065843	0.688075

	median_house_value
longitude	-0.045967
latitude	-0.144160
housing_median_age	0.105623
total_rooms	0.134153
total_bedrooms	0.049686
population	-0.024650
households	0.065843
median_income	0.688075
median_house_value	1.000000

```
[23]: ##### 12. Group the data and sort in ascending order
df.groupby(by='median_house_value').count().sort_values('median_house_value',
↪ascending=False).head(10)
```


```
[23]:
```

	longitude	latitude	housing_median_age	total_rooms	\
median_house_value					
500001.0	965	965	965	965	
500000.0	27	27	27	27	
499100.0	1	1	1	1	
499000.0	1	1	1	1	
498800.0	1	1	1	1	
498700.0	1	1	1	1	
498600.0	1	1	1	1	
498400.0	1	1	1	1	
497600.0	1	1	1	1	
497400.0	1	1	1	1	

	total_bedrooms	population	households	median_income	\
median_house_value					
500001.0	958	965	965	965	
500000.0	27	27	27	27	
499100.0	1	1	1	1	
499000.0	1	1	1	1	
498800.0	1	1	1	1	
498700.0	1	1	1	1	
498600.0	1	1	1	1	
498400.0	1	1	1	1	
497600.0	1	1	1	1	


497400.0	1	1	1	1
ocean_proximity				
median_house_value				
500001.0	965			
500000.0	27			
499100.0	1			
499000.0	1			
498800.0	1			
498700.0	1			
498600.0	1			
498400.0	1			
497600.0	1			
497400.0	1			

```
[25]: ##### 13. Plot the data using scatterplot
_ = df.plot.scatter('median_income', 'median_house_value')
```




```
[26]: _ = sns.boxplot(y='median_house_value', data=df)
```






```
[27]: _ = sns.violinplot(y='median_house_value', data=df)
```




```
[28]: _ = sns.jointplot(x="median_income", y="median_house_value", data=df,   
kind="reg")
```




```
[30]: ##### 14. Removing outliers  
df.drop(df[df['median_house_value']>500000].index, inplace=True)
```

```
[31]: _ = sns.violinplot(y='median_house_value', data=df)
```



```
[32]: _ = sns.jointplot(x="median_income", y="median_house_value", data=df,   
    ↪ kind="reg")
```



### Preparing the data for training and testing

1. Divide our independent and dependent variable into two separate variables.
2. Split the data into training and testing datasets.

[42]: # 1)

```
X = df.iloc[:,8].values.reshape(-1,1) # input
y = df.iloc[:,7].values # output (dependent variable)
```

[43]: # 2) Splitting our data into training and testing sets

```
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪shuffle=False, random_state=SEED)
```

### Train the Model

```
[44]: # Import the linear regression algorithm
from sklearn.linear_model import LinearRegression

regressor = LinearRegression()


# Train the model
regressor.fit(X_train, y_train)
```

[44]: LinearRegression()

### 15. Visualize the trained model

```
[45]: #  $y=mx+c$  (Linear regression model)
line = regressor.coef_*X + regressor.intercept_
```

```
[46]: # Lets plot this on the scatter plot
plt.scatter(X,y)
plt.plot(X, line, 'r')
plt.xlabel("Median House Value")
plt.ylabel("Median Income")
plt.title("Median House Value vs Median Income")
plt.show()
```



## 1.6 Rubrics:

Criteria	Scoring
Data Preprocessing	20 pts.
Training Performance	20 pts.
Model Evaluation Metrics	20 pts.
Visualization	20 pts.
Explanation & Analysis	20 pts.

[ ]: